

EX.NO:04

14.08.25

# BUILD A SIMPLE FEED FORWARD NEURAL NETWORK TO RECOGNIZE HANDWRITTEN CHARACTER

Aim:

To build a simple feed forward neural network to recognize handwritten character.

Objective:

1. To load and preprocess the mnist dataset for neural network input
2. To build feed forward neural network with hidden layers
3. To train the model using stochastic gradient descent optimizer and sparse categorical cross-entropy loss
4. Evaluate the trained model on test data and measure its accuracy.
5. To predict the class of given handwritten image.

Pseudocode

START

Load MNIST dataset (training and testing data)

Flatten each image from 28x28 to 784 features

Normalize pixel values to range [0,1]

Create a sequential neural network:

Layer 1: Dense (128 neurons, ReLU activation)

Layer 2: Dense (64 neurons, ReLU activation)

Output layer: Dense (10 neurons, softmax)



## Eg Training

epoch	Accuracy	Loss
1	0.9929	0.0232
2	0.9968	0.0128
3	0.9976	0.0099
4	0.9976	0.0088
5	0.9987	0.0058

## Training

Epoch	Loss
1	1.0594
2	0.3860
3	0.3290
4	0.2974
5	0.2726

Test Accuracy: 92.82%



activation).

Compile model

optimizer = stochastic gradient descent

loss = sparse categorical\_crossentropy

Metric = accuracy

Train model on training data for 5 epochs

Evaluate model on testing data

print test accuracy

END

observation

The loss decreases with each epoch, showing that the model is learning.

Accuracy improves steadily during training

Testing

overall accuracy of the model on the entire dataset = ~~0.9828~~ 92.82%

Results

Successfully build a simple feed forward neural network to recognize handwritten character.

~~2/11/21~~

The screenshot displays a Jupyter Notebook interface. On the left, a file explorer shows a directory named 'DEEP LEARNING' containing several files, including 'L4.2.ipynb' which is selected. The notebook content on the right shows a series of code cells for setting up a PyTorch environment and defining a simple neural network.

```
[37]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
from torchvision import transforms
from torch.utils.data import DataLoader

[38]: transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

[39]: train_dataset = torchvision.datasets.MNIST(
    root='./data', train=True, transform=transform, download=True
)
test_dataset = torchvision.datasets.MNIST(
    root='./data', train=False, transform=transform, download=True
)

[40]: train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)

[41]: import torch.nn as nn
class SimpleFFN(nn.Module):
    def __init__(self):
        super(SimpleFFN, self).__init__()
        self.flatten = nn.Flatten()
        self.fc1 = nn.Linear(28*28, 128)
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, 10)
        self.relu = nn.ReLU()
```



L4.2.ipynb (4) - JupyterLab

IndentationError fix

What is PyTorch? | Data Science

difference between tensorflow

10.1.38.19/user/ra2311047010011/lab/tree/DEEP%20LEARNING/L4.2.ipynb

Verify it's you

FileEditViewRunKernelTabsSettingsHelp

Filter files by name

/ DEEP LEARNING /

Name	Last Modified
data	7 days ago
Breast_canc...	14 days ago
Exp2.ipynb	14 days ago
installpytor...	30 minutes ago
L_3.ipynb	7 days ago
L1.1.ipynb	14 days ago
L1.ipynb	14 days ago
L2.ipynb	7 days ago
L3.ipynb	7 days ago
L4.2.ipynb	2 minutes ago
L4.ipynb	30 minutes ago
Untitled.ipy...	14 days ago

L4.ipynbinstallpytorch.ipynbL4.2.ipynb

Code

Python 3 (ipykernel)

```
self.fc3 = nn.Linear(64, 10)
self.relu = nn.ReLU()

def forward(self, x):
    x = self.flatten(x)
    x = self.relu(self.fc1(x))
    x = self.relu(self.fc2(x))
    x = self.fc3(x)
    return x

[42]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
      model = SimpleFFN().to(device)

[43]: criterion = nn.CrossEntropyLoss()
      optimizer = optim.SGD(model.parameters(), lr=0.01)
      epochs = 5

[44]: for epoch in range(epochs):
      running_loss = 0.0
      for images, labels in train_loader:
          images, labels = images.to(device), labels.to(device)

          outputs = model(images)
          loss = criterion(outputs, labels)


          optimizer.zero_grad()
          loss.backward()
          optimizer.step()

          running_loss += loss.item()
```

Simple03Python 3 (ipykernel) | IdleMem: 1.31 GB

Mode: CommandLn 1, Col 1L4.2.ipynb0

ENG09:1622-08-2025



```
running_loss += loss.item()

print(f"Epoch [{epoch+1}/{epochs}], Loss: {running_loss/len(train_loader):.4f}")

Epoch [1/5], Loss: 1.0231
Epoch [2/5], Loss: 0.3828
Epoch [3/5], Loss: 0.3246
Epoch [4/5], Loss: 0.2925
Epoch [5/5], Loss: 0.2678

[45]: print("Training Finished")
Training Finished

[46]: correct, total = 0, 0
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f"Test Accuracy: {100 * correct / total:.2f}%")

Test Accuracy: 92.41%
```